

CSCE 463/612 Networks and Distributed Processing Fall 2020

Transport Layer V

Dmitri Loguinov
Texas A&M University

October 15, 2020

Chapter 3: Roadmap

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP

- Segment structure
- Reliable data transfer
- Flow control
- Connection management

3.6 Principles of congestion control

3.7 TCP congestion control

Principles of Congestion Control



Congestion:

- Informally: “too many sources sending too much data too fast for the *network* to handle”
- Different from flow control!
- Manifestations:
 - Lost packets (buffer overflows)
 - Delays (queueing in routers)
- Important networking problem



Causes/Costs of Congestion: Scenario 1


- Two senders, two receivers
- One router of capacity C , infinite buffers, no loss
- No retransmission



Cost 1: queuing delays in congested routers




Causes/Costs of Congestion: Scenario 2

- One router, *finite* buffers (pkt loss is possible now)
- Sender retransmission of lost packet
- During congestion $2\lambda_{net} = 2(\lambda_{in} + \lambda_{retx}) = C$



Causes/Costs of Congestion: Scenario 2

- We call $\lambda_{in} = \lambda_{out}$ **goodput** and λ_{net} **throughput**
 - Case A: pkts never lost while $\lambda_{net} < C/2$ (not realistic)
 - Case B: pkts are lost when λ_{net} is “sufficiently large,” but timeouts are perfectly accurate (not realistic either)
 - Case C: same as B, but timer is not perfect (duplicate packets are possible)




Cost 2: retransmission of lost packets and premature timeouts increase network load, reduce *flow's own* goodput

Causes/Costs of Congestion: Scenario 3

- Multihop case
 - Timeout/retransmit
 - $R2 = 50 \text{ Mbps}$, $R1 = R3 = R4 = 100 \text{ Mbps}$
 - Flow C-A: sends 90 Mbps

Cost 3: congestion causes goodput reduction for *other* flows



Approaches Towards Congestion Control

Two broad approaches towards congestion control:

End-to-end:

- No **explicit** feedback from network
- Congestion ***inferred*** by end-systems from observed loss/delay
 - Approach taken by TCP (relies on loss)

ATM = Asynchronous
Transfer Mode


Network-assisted:

- Routers provide feedback to end systems
 - Single bit indicating congestion (DECbit, TCP/IP ECN)
 - Two bits (ATM)
 - Explicit rate senders should send at (ATM)

Case Study: ATM ABR Congestion Control

- For network-assisted protocols, the logic can be **binary**:
 - Path underloaded, increase rate
 - Path congested, reduce rate
- It can also be **ternary**
 - Increase, decrease, hold steady
 - ATM ABR (Available Bit Rate) profile
- **RM (resource management) packets (cells):**
 - Sent by sender, interspersed with data cells
 - Bits in RM cell set by switches/routers
 - **NI bit**: no increase in rate (impending congestion)
 - **CI bit**: reduce rate (congestion in progress)
 - RM cells returned to sender by receiver, with bits intact

Case Study: ATM ABR Congestion Control



- Additional approach is to use a two-byte ER (explicit rate) field in RM cell
 - Congested switch may lower ER value
 - Senders obtain the maximum supported rate on their path
- Issues with network-assisted congestion control?

Chapter 3: Roadmap

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP

- Segment structure
- Reliable data transfer
- Flow control
- Connection management

3.6 Principles of congestion control

3.7 TCP congestion control

TCP Congestion Control

- TCP congestion control has a variety of algorithms developed over the years
 - **TCP Tahoe** (1988), **TCP Reno** (1990), TCP SACK (1992)
 - TCP Vegas (1994), TCP New Reno (1996)
 - High-Speed TCP (2002), Scalable TCP (2002)
 - FAST TCP (2004), TCP Illinois (2006)
- Many others: H-TCP, CUBIC TCP, L-TCP, TCP Westwood, TCP Veno (Vegas + Reno), TCP Africa
- Linux: BIC TCP (2004), CUBIC TCP (2008)
- Vista and later: Compound TCP (2005)
 - Server 2019 switched to CUBIC
- Google: BBR (2016)

TCP Congestion Control


- **End-to-end** control (no network assistance)
- Sender limits transmission:
$$\text{LastByteSent} - \text{LastByteAked} \leq \text{CongWin}$$
- CongWin is a function of perceived network congestion
- The *effective* window is the minimum of CongWin, flow-control window carried in the ACKs, and sender's own buffer space
- How does sender perceive congestion?
 - Loss event = timeout or 3 duplicate acks
- TCP sender reduces rate (CongWin) after loss event
- Three mechanisms:
 - AIMD (congestion avoidance)
 - Slow start
 - Conservative after timeout events

TCP AIMD (Additive Increase, Multiplicative Decrease)

Additive increase: increase CongWin by 1 MSS every RTT in the absence of loss events: *probing*

Multiplicative decrease: cut CongWin in half after fast retransmit (3-dup ACKs)

Peaks are different: # of flows or RTT changes



TCP Equations

- To better understand TCP, we next examine its AIMD equations (**congestion avoidance**)
- Assume that W is the window size in pkts and $B = \text{CongWin}$ is the same in bytes ($B = \text{MSS} * W$)
- General form (loss detected through 3-dup ACK):

$$W = \begin{cases} W + \frac{1}{W} & \text{per ACK} \\ W/2 & \text{per loss} \end{cases}$$

- Reasoning
 - For each window of size W , we get exactly W acknowledgments in one RTT (assuming no loss!)
 - This increases window size by “roughly” 1 packet per RTT

TCP Equations

$$W = \begin{cases} W + \frac{1}{W} & \text{per ACK} \\ W/2 & \text{per loss} \end{cases}$$

- What is the equation in terms of $B = MSS * W$?

$$B = \begin{cases} B + \frac{MSS^2}{B} & \text{per ACK} \\ B/2 & \text{per loss} \end{cases}$$

- Equivalently, TCP increases B by MSS per RTT
- What is the rate of TCP given that its window size is B (or W)?
- Since TCP sends a full window of pkts per RTT, its ideal rate can be written as:


$$r = \frac{B}{RTT + L/R} \approx \frac{B}{RTT} = \frac{MSS * W}{RTT}$$

TCP Slow Start

- When connection begins, $\text{CongWin} = 1 \text{ MSS}$
 - Example: $\text{MSS} = 500 \text{ bytes}$ and $\text{RTT} = 200 \text{ msec}$
 - Q: initial rate?
- A: 20 Kbits/s
- Available bandwidth may be much larger than MSS/RTT
 - Desirable to quickly ramp up to a “respectable” rate
- Solution: **Slow Start (SS)**
 - When a connection begins, it increases rate exponentially fast until first loss or receiver window is reached
 - Term “slow” is used to distinguish this algorithm from earlier TCPs which directly jumped to some huge rate

TCP Slow Start (More)

- Slow start
 - Double CongWin every RTT
- Done by incrementing CongWin for every ACK received:
 - $W = W + 1$ per ACK
(or $B = B + MSS$)
- Summary: initial rate is slow but ramps up exponentially fast




Refinement

- **TCP Tahoe** responds only to timeouts:
 - $\text{Threshold} = \text{CongWin}/2$
 - CongWin is set to 1 MSS
 - Slow start until threshold is reached; then move to AIMD congestion avoidance
- **TCP Reno** loss:
 - Timeout: same as Tahoe
 - 3 dup ACKs: CongWin is cut in half (original idea was called **fast recovery**, now part of AIMD)

loss detected via triple dup ACK

previous timeout



Fast Recovery Philosophy:

Three dup ACKs indicate that network is capable of delivering subsequent segments

Timeout before 3-dup ACK is more alarming

Refinement (More)

- Initial slow start ends when either
 - Loss occurs
 - Initial threshold is reached
- **Initial threshold** is usually set to the receiver's advertised window

Implementation:

- Variable `ssthresh` is the “slow start threshold”
- At loss events, `ssthresh` is set to $\text{CongWin} / 2$

TCP Reno Sender Congestion Control

Event	State	TCP Sender Action	Commentary
ACK receipt for previously unacked data	Slow Start (SS)	CongWin += MSS, If (CongWin >= ssthresh) { Set state to "Congestion Avoidance" }	Results in a doubling of CongWin every RTT
ACK receipt for previously unacked data	Congestion Avoidance (CA)	CongWin += $MSS^2 / \text{CongWin}$	Additive increase, resulting in increase of CongWin by 1 MSS every RTT
Loss event detected by triple duplicate ACK	SS or CA	ssthresh = max(CongWin/2, MSS) CongWin = ssthresh Set state to "Congestion Avoidance"	Fast recovery, implementing multiplicative decrease
Timeout	SS or CA	ssthresh = max(CongWin/2, MSS) CongWin = MSS Set state to "Slow Start"	Enter slow start
Duplicate ACK	SS or CA	Increment duplicate ACK count for segment being acked	CongWin and Threshold not changed

TCP Congestion Control

- Summary of TCP Reno:

